# Doing More with Data: pandas

Introduction to Python

Data Sciences Institute, University of Toronto Instructor: A Mahfouz | TA: Kaylie Lau July 2022

# Contents:

- 1. Setup
- 2. Intro to pandas
- 3. Getting data
- 4. Profiling and initial data exploration: changing data types, descriptive statistics
- 5. Wrangling and plotting: concatenating, merging, adding and removing columns, filtering and selecting, null values, grouping and aggregating, plotting
- 6. Writing to file
- 7. More wrangling: reshaping, applying functions

# Setup

### Anaconda

Install pandas if you haven't already. If you're not sure, go through the steps below.

- 1. Open Anaconda Prompt. You may have to do this in admin mode.
- 2. Type conda install pandas and press enter to run.
- 3. Press enter to answer Y to any prompts.

We will also need openpyx1, which can be installed by running the command conda install openpyx1 in Anaconda Prompt.

# Google Colab

If you're using Colab, you're all set! pandas and openpyx1 are already installed.

# Data

This module uses four datasets: <u>bike thefts</u>, <u>TTC subway delays and subway delay reason</u> <u>codes</u>, and <u>neighbourhood profiles</u>. All four are available in the course repo, and originally come from Toronto Open Data.

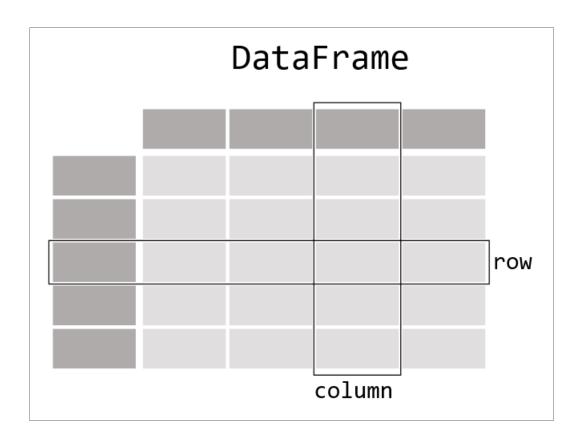
# pandas

## What is pandas?

pandas is a package for data analysis and manipulation. (The name is a reference to panel data, not the animal.) It gives us data frames, which represent data in a table of columns and rows, and functions to manipulate and plot them. pandas also provides a slew of functions for reading and writing data to a variety of sources, including files, SQL databases, and compressed binary formats.

### DataFrames

Columns are labeled with their names. Rows also have a label, or *index*. If row labels are not specified, pandas uses numbers as the default. Each column is a *Series*, or one-dimensional array, where values share a data type. Unlike numpy arrays, DataFrames can have columns of different data types. However, like arrays and lists, **DataFrames are mutable** -- this means that if more than one variable refers to the same DataFrame, updating one updates them all!



# Getting data

We can create a DataFrame manually with DataFrame() constructor. If a dictionary is passed to DataFrame(), the keys become column names, and the values become the rows. Calling just DataFrame() creates an empty DataFrame to which data can be added later.

#### Out[2]:

	name	avg_lifespan	quantity
0	sugar maple	300	53
1	black oak	100	207
2	white ash	260	178
3	douglas fir	450	93

We can create an individual column with Series(). The name argument corresponds to a column name.

#### Out[3]:

- 0 deciduous
- 1 deciduous
- 2 deciduous
- 3 evergreen

Name: foliage, dtype: object

#### Data from csv

Of course, we're more likely to load data into a DataFrame than to create DataFrames manually. pandas has read functions for different file formats. To read data from a csv or other delimited file, we use pd.read\_csv(), then pass in the local file path or the URL of the csv to read. pandas will infer the data type of each column based on the values in the first chunk of the file loaded.

```
In [4]: thefts = pd.read_csv('../data/bicycle-thefts - 4326.csv')
```

## Profiling and initial data cleaning

We got our data, but now we need to understand what's in it. We can start to understand the DataFrame by checking out its dtypes and shape attributes, which give column data types and row by column dimensions, respectively. Note that object is pandas 'way of saying values are represented as string data.

```
In [5]: thefts.shape
Out[5]:
  (25569, 33)
```

## In [6]: thefts.dtypes

### Out[6]:

_id	int64
OBJECTID	int64
event_unique_id	object
Primary_Offence	object
Occurrence_Date	object
Occurrence_Year	int64
Occurrence_Month	object
Occurrence_DayOfWeek	object
Occurrence_DayOfMonth	int64
Occurrence_DayOfYear	int64
Occurrence_Hour	int64
Report_Date	object
Report_Year	int64
Report_Month	object
Report_DayOfWeek	object
Report_DayOfMonth	int64
Report_DayOfYear	int64
Report_Hour	int64
Division	object
City	object
Hood_ID	object
NeighbourhoodName	object
Location_Type	object
Premises_Type	object
Bike_Make	object
Bike_Model	object
Bike_Type	object
Bike_Speed	int64
Bike_Colour	object
Cost_of_Bike	float64
Status	object
ObjectId2	int64
geometry	object
dtype: object	

head() s and tail() s

To check out the first few rows, we can call the DataFrame head() method. Similarly, we can see the last few rows with the tail() method. Five rows are shown by default, but we can change that by passing an integer as an argument.

<pre>In [7]: thefts.head()</pre>						
Out[7]:	OBJECTID	event_unique_id	Primary_Offence	Occurrence_Date	Occurrence_Year	Occurrence_Mont
0 1	17744	GO-20179016397	THEFT UNDER	2017-10-03T00:00:00	2017	Octobe
1 2	17759	GO-20172033056	THEFT UNDER - BICYCLE	2017-11-08T00:00:00	2017	Novembe
<b>2</b> 3	17906	GO-20189030822	THEFT UNDER - BICYCLE	2018-09-14T00:00:00	2018	Septembe
3 4	17962	GO-2015804467	THEFT UNDER	2015-05-07T00:00:00	2015	Ma
4 5	17963	GO-20159002781	THEFT UNDER	2015-05-16T00:00:00	2015	Ma

UNDER

In [8]: # Last 3
 thefts.tail(3)

Out[8]:

	_id	OBJECTID	event_unique_id	Primary_Offence	Occurrence_Date	Occurrence_Year	Occurre
25566	25567	11462	GO-20169005434	THEFT UNDER	2016-06-04T00:00:00	2016	
25567	25568	11695	GO-20161170896	THEFT UNDER	2016-07-04T00:00:00	2016	
25568	25569	11883	GO-20169007653	THEFT UNDER - BICYCLE	2016-07-22T00:00:00	2016	

### Renaming columns

Most, but not all, of the bike theft columns follow the same naming convention. For convenience's sake, though, let's convert the column names to all lowercase. We can do this with the DataFrame rename() method. rename() accepts either a dictionary with current column names as the keys and new names as the values, or the name of a function to transform names. Let's write a function.

Let's also rename cost\_of\_bike so it follows the pattern of the other bike attribute columns.

```
In [10]: thefts = thefts.rename(columns={'cost_of_bike':'bike_cost'})
# view column names
print(list(thefts))
```

['\_id', 'objectid', 'event\_unique\_id', 'primary\_offence', 'occurrence\_date', 'occurrence\_year', 'occurrence\_month', 'occurrence\_dayofweek', 'occurrence\_dayofmonth', 'occurrence\_dayofyear', 'occurrence\_hour', 'report\_date', 'report\_year', 'report\_month', 'report\_dayofweek', 'report\_dayofmonth', 'report\_dayofyear', 'report\_hour', 'division', 'city', 'hood\_id', 'neighbourhoodname', 'location\_type', 'premises\_type', 'bike\_make', 'bike\_model', 'bike\_type', 'bike\_speed', 'bike\_colour', 'bike\_cost', 'status', 'objectid2', 'geometry']

### Profiling columns

It can be useful to focus on a subset of columns, particularly to understand value sets. To select a single column in a DataFrame, we can supply the name of the column in square brackets, just like we did when accessing values in a dictionary. pandas will return the column as a Series. To get unique values, we can use the unique() Series method. If we want to count how many times each value appears, we can use the value\_counts() method.

```
In [11]: thefts['status']
Out[11]:
0
             STOLEN
1
         RECOVERED
2
             STOLEN
3
             STOLEN
             STOLEN
            . . .
25564
             STOLEN
25565
            STOLEN
25566
            STOLEN
25567
            STOLEN
25568
             STOLEN
Name: status, Length: 25569, dtype: object
In [12]: thefts['status'].unique()
Out[12]:
array(['STOLEN', 'RECOVERED', 'UNKNOWN'], dtype=object)
In [13]: thefts['status'].value_counts()
Out[13]:
STOLEN
              24807
                454
UNKNOWN
RECOVERED
                308
```

Name: status, dtype: int64

We can summarize numeric Series much like we did with numpy functions.

```
In [14]: thefts['bike_cost'].median()

Out[14]:
600.0

In [15]: thefts['bike_cost'].quantile(0.9)

Out[15]:
2000.0
```

## info()

We can get an overview of the DataFrame by profiling it with the info() method. info() prints a lot of information about a DataFrame, including:

- the shape as the number of rows and columns
- column names and their dtype
- the number of non-null values in each column
- how big the DataFrame is in terms of memory usage

The bicycle theft data looks quite complete, though some records are missing bike descriptors like bike\_make, bike\_model, bike\_colour, and bike\_cost. Most of the column dtypes make sense. We'll want to convert the dates to proper dates. We may also want to convert string columns with limited value sets, like status, to categorical data.

In [16]: thefts.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 25569 entries, 0 to 25568
Data columns (total 33 columns):

Data	columns (total 33 colu		
#	Column	Non-Null Count	Dtype
0	_id	25569 non-null	int64
1	objectid	25569 non-null	int64
2	event_unique_id	25569 non-null	object
3	primary_offence	25569 non-null	object
4	occurrence_date	25569 non-null	object
5	occurrence_year	25569 non-null	int64
6	occurrence_month	25569 non-null	object
7	occurrence_dayofweek	25569 non-null	object
8	occurrence_dayofmonth	25569 non-null	int64
9	occurrence_dayofyear	25569 non-null	int64
10	occurrence_hour	25569 non-null	int64
11	report_date	25569 non-null	object
12	report_year	25569 non-null	int64
13	report_month	25569 non-null	object
14	report_dayofweek	25569 non-null	object
15	report_dayofmonth	25569 non-null	int64
16	report_dayofyear	25569 non-null	int64
17	report_hour	25569 non-null	int64
18	division	25569 non-null	object
19	city	25569 non-null	object
20	hood_id	25569 non-null	object
21	neighbourhoodname	25569 non-null	object
22	location_type	25569 non-null	object
23	premises_type	25569 non-null	object
24	bike_make	25448 non-null	object
25	bike_model	15923 non-null	object
26	bike_type	25569 non-null	object
27	bike_speed	25569 non-null	int64
28	bike_colour	23508 non-null	object
29	bike_cost	23825 non-null	float64
30	status	25569 non-null	object
31	objectid2	25569 non-null	int64
32	geometry	25569 non-null	object

dtypes: float64(1), int64(12), object(20)
memory usage: 6.4+ MB

### Changing data types

Before exploring the bike theft data further, let's fix the date and categorical columns. To convert a column to datetime, we use the pd.to\_datetime() function, passing in the column to convert, and reassign the output back to the column we're converting. pandas knows how to convert the dates in the bike thefts data, but for less common formats, it is necessary to use the format keyword argument to specify how dates should be parsed. format strings use strftime codes. See <a href="https://strftime.org/">https://strftime.org/</a> for a cheat sheet.

```
Out[17]:
        2017-10-03
0
1
        2017-11-08
2
        2018-09-14
3
        2015-05-07
4
        2015-05-16
           . . .
25564
        2015-04-01
25565
        2016-05-16
25566
        2016-06-04
25567
        2016-07-04
25568
        2016-07-22
Name: occurrence_date, Length: 25569, dtype: datetime64[ns]
```

```
Out[18]:
0
        2017-10-03
1
        2017-11-08
2
        2018-09-17
3
        2015-05-14
4
        2015-05-16
           . . .
25564
        2015-04-01
25565
        2016-05-16
25566
        2016-06-07
25567
        2016-07-04
25568
        2016-07-23
Name: report_date, Length: 25569, dtype: datetime64[ns]
```

All other data type conversions can be done with the astype() method. If we were converting to a number, pd.to\_numeric() provides an easy way to convert without having to pick a specific numeric data type.

```
Out[19]:
0
             STOLEN
1
         RECOVERED
2
             STOLEN
3
             STOLEN
4
             STOLEN
25564
             STOLEN
25565
             STOLEN
25566
             STOLEN
25567
             STOLEN
25568
             STOLEN
Name: status, Length: 25569, dtype: category
Categories (3, object): ['RECOVERED', 'STOLEN', 'UNKNOWN']
```

We can select and convert multiple columns at once by passing a list of columns in the square brackets., then using <code>.astype()</code>.

```
In [20]: thefts[['location_type', 'premises_type']] = thefts[['location_type', 'premises_type']].astype('category')
# check data types
thefts[['location_type', 'premises_type']].dtypes
```

#### Out[20]:

location\_type category premises\_type category

dtype: object

### describe()

To get a sense of the values in a DataFrame, we can use the describe() method. describe() summarizes only numeric columns by default. Passing the include='all' argument will produce summary statistics for other columns as well.

Oı	ut	[	21	.]	

						ouc[zi].
occurre	occurrence_date	primary_offence	event_unique_id	objectid	_id	
25569.	25569	25569	25569	25569.000000	25569.000000	count
	NaN	66	22771	NaN	NaN	unique
	NaN	THEFT UNDER	GO-20201550944	NaN	NaN	top
	NaN	11904	14	NaN	NaN	freq
2017.	2017-09-04 03:39:28.321013504	NaN	NaN	12909.173218	12785.000000	mean
2009.	2009-09-01 00:00:00	NaN	NaN	1.000000	1.000000	min
2016.	2016-01-06 00:00:00	NaN	NaN	6456.000000	6393.000000	25%
2017.	2017-09-05 00:00:00	NaN	NaN	12918.000000	12785.000000	50%
2019.	2019-06-20 00:00:00	NaN	NaN	19360.000000	19177.000000	75%
2020.	2020-12-30 00:00:00	NaN	NaN	25806.000000	25569.000000	max

	_id	objectid	event_unique_id	primary_offence	occurrence_date	occurre
std	7381.278853	7448.318562	NaN	NaN	NaN	1.9

# Practice

- 1. Explore unique valuesets for some of the other columns in thefts. Identify at least one that may make sense as categorical variables and convert it.
- 2. Count values for location\_type. Where did most reported thefts happen?
- 3. Find the 10th percentile of bike\_cost.

Wrangling and Plotting

### Combining datasets: concatenation

Just as pandas has read\_csv() for flat files, there is a read\_excel() function to load Excel files.

The TTC publishes subway delay data as a multi-sheet Excel workbook, with a month's worth of data per sheet. read\_excel() loads just the first sheet in an Excel file by default. To load all sheets, pass in the keyword argument sheet\_name=None. The result is a dictionary, where each key is the sheet name and each value is a DataFrame with the contents of the sheet.

To combine them all, we create an empty DataFrame, then loop through the dictionary items and use pd.concat() to append data. concat() takes a list of DataFrames to combine. Since we did not specify an index, row labels are numbers: the first row of each sheet has an index of 0, and so on. To reset row labels so that they are sequential again, we set ignore\_index=True.

```
Adding 1216 rows from January21
Adding 1245 rows from Feb 21
Adding 1167 rows from March '21
Adding 1170 rows from April '21
Adding 1168 rows from May '21
Adding 1265 rows from June 21
Adding 1244 rows from July 21
Adding 1273 rows from August 21
Adding 1433 rows from Sept 21
Adding 1560 rows from Oct 21
Adding 1771 rows from Nov 21
Adding 1858 rows from December21
Out[25]:
(16370, 10)
```

In [26]: all\_delays.head()

### Out[26]:

	Date	Time	Day	Station	Code	Min Delay	Min Gap	Bound	Line	Vehicle
0	2021-01-01	00:33	Friday	<b>BLOOR STATION</b>	MUPAA	0	0	N	YU	6046
1	2021-01-01	00:39	Friday	SHERBOURNE STATION	EUCO	5	9	Е	BD	5250
2	2021-01-01	01:07	Friday	KENNEDY BD STATION	EUCD	5	9	E	BD	5249
3	2021-01-01	01:41	Friday	ST CLAIR STATION	MUIS	0	0	NaN	YU	0
4	2021-01-01	02:04	Friday	SHEPPARD WEST STATION	MUIS	0	0	NaN	YU	0

# Practice

- 1. The TTC delays data includes a reason code for the delay. Code definitions, however, are in a separate Excel file, ttc-subway-delay-codes.xlsx. This file has been modified slightly to make it easier to work with. Codes are split between two tabs. Load both to a DataFrame, delay\_reasons.
- 2. Rename the columns in both all\_delays and delay\_reasons so that all letters are lowercase. **Challenge**: rename the columns to replace spaces with underscores as well as convert all letters to lowercase.
- 3. Explore all\_delays. Do any columns have missing values? What was the median Min Delay? How many unique lines are in the data?

### Combining datasets: merging

Ideally, the delays data would include code descriptions. We can get descriptions into all\_delays by *merging* in delay\_reasons. Merging is analagous to joining in SQL databases. To merge two DataFrames, we pass them as arguments to the pd.merge(). Then, we specify how to merge the two DataFrames and what column names to merge on.

Let's review the all\_delays and delay\_reasons DataFrames. code is equivalent to rmenu\_code. If we pass in all\_delays as the first DataFrame, then it will be the left frame, and delay\_reasons the right one. We want to keep all the delay records, even if there isn't a matching code in delay\_reasons, so we will perform a left join.

In [29]: all\_delays.head(2)

#### Out[29]:

	date	time	day	station	code	min_delay	min_gap	bound	line	vehicle
0	2021-01-01	00:33	Friday	BLOOR STATION	MUPAA	0	0	N	YU	6046
1	2021-01-01	00:39	Friday	SHERBOURNE STATION	EUCO	5	9	Е	BD	5250

In [30]: delay\_reasons.head(2)

### Out[30]:

	rmenu_code	code_description	sub_or_srt		
0	EUAC	Air Conditioning	SUB		
1	EUAL	Alternating Current	SUB		

Out[31]:

	date	time	day	station	code	min_delay	min_gap	bound	line	vehicle	rmenu_co
0	2021-01-01	00:33	Friday	BLOOR STATION	MUPAA	0	0	N	YU	6046	MUPA
1	2021-01-01	00:39	Friday	SHERBOURNE STATION	EUCO	5	9	E	BD	5250	EUC
2	2021-01-01	01:07	Friday	KENNEDY BD STATION	EUCD	5	9	E	BD	5249	EUC

# Practice

Try performing an inner join with all\_delays and delay\_reasons. Assign the result to a different variable name. Compare the shapes of delays\_w\_reasons and this new DataFrame. Did we get the same number of rows with the inner join?

### drop()

The resulting DataFrame has both our join columns, which is redundant. We can drop one with the drop() DataFrame method, passing in the column name(s) we want to drop in the columns keyword argument.

station

#### Out[32]:

date

time

day

Pas Ass Act No	6046	YU	N	0	0	MUPAA	BLOOR STATION	Friday	00:33	2021-01-01	0
C	5250	BD	Е	9	5	EUCO	SHERBOURNE STATION	Friday	00:39	2021-01-01	1
Consec Del Dela	5249	BD	E	9	5	EUCD	KENNEDY BD Station	Friday	01:07	2021-01-01	2

code min\_delay min\_gap bound line vehicle

code de

### Creating new columns

Adding a column to a DataFrame looks like adding a key-value pair to a dictionary. At its simplest, we can assign a single value to repeat down a column.

```
In [33]: delays_w_reasons['year'] = 2021
     delays_w_reasons['year'].unique()

Out[33]:
array([2021], dtype=int64)
```

We can also write an expression and store the resulting values in a new column.

#### Out[34]:

	min_delay	hour_delay
0	0	0.00
1	5	0.08
2	5	0.08
3	0	0.00
4	0	0.00

It is also possible to extract parts of a datetime column with the dt accessor.

```
Out[35]:
0
           1
1
           1
2
           1
3
           1
           1
16365
         12
16366
         12
16367
         12
16368
         12
         12
16369
Name: month, Length: 16370, dtype: int64
```

# Practice

- Create a new integer column, hour, that contains the hour in which a delay occurred.

  Hint: Use pd.to\_datetime(), specifying that the format is hour:minutes.
- Drop the year column.

### Filtering and selecting data

Let's take another look at the TTC subway delay data. There are only 4 subway lines in Toronto, but describe() reported 17 unique values.

Looks like some of the line values should be updated (YU/BD variants) and others should be dropped (e.g., 36 FINCH WEST, NaNs). Luckily there don't seem to be too many affected records, though the NaNs are not shown.

```
In [38]: delays_w_reasons['line'].value_counts()
```

#### Out[38]: YU 8880 BD5734 SHP 657 SRT 656 YU/BD 346 YUS 18 YU / BD 17 YU & BD LINES 1 41 KEELE 1 52 1 35 JANE 1 999 1 YUS & BD 1 36 FINCH WEST 1 SHEP 1 YONGE/UNIVERSITY/BLOOR 1 YUS/BD 1 Name: line, dtype: int64

```
.loc[] and isna()
```

To find the records with no line, we can use .loc[], which lets us access rows and columns with either a boolean array or row/column labels.

In this case, the boolean array is the product of the <code>isna()</code> Series method.

Out[39]:

Out[39]	, •									
	date	time	day	station	code	min_delay	min_gap	bound	line	vehi
495	2021-01-13	15:22	Wednesday	FINCH WEST STATION	MUSAN	3	6	S	NaN	57.
513	2021-01-13	22:08	Wednesday	EGLINTON WEST STATION	PUMEL	0	0	NaN	NaN	
1044	2021-01-27	22:00	Wednesday	YONGE- UNIVERSITY AND B	MUO	0	0	NaN	NaN	
1045	2021-01-27	23:00	Wednesday	FINCH STATION	MUO	0	0	NaN	NaN	
1362	2021-02-04	01:45	Thursday	LAWRENCE STATION	TUSC	0	0	S	NaN	55
1679	2021-02-11	01:12	Thursday	GREENWOOD CARHOUSE	MUIE	0	0	NaN	NaN	
2179	2021-02-22	08:27	Monday	BICHMOUNT DIVISION	MUIE	0	0	NaN	NaN	
2204	2021-02-22	22:33	Monday	BLOOR STATION	SUAP	4	9	N	NaN	60
2206	2021-02-22	23:36	Monday	EGLINTON	MUO	0	0	NaN	NaN	

	date	time	day	station	code	min_delay	min_gap	bound	line	vehi
				STATION						
3039	2021-03-17	05:15	Wednesday	INGLIS BUILDING	PUMEL	0	0	NaN	NaN	
3330	2021-03-24	19:13	Wednesday	INGLIS BUILDING	PUMEL	0	0	NaN	NaN	
3407	2021-03-26	09:03	Friday	WILSON YARD (SOUTH TAI	PUTO	0	0	NaN	NaN	
3557	2021-03-30	00:36	Tuesday	INGLIS BUILDING	PUMEL	0	0	NaN	NaN	
3944	2021-04-08	23:45	Thursday	DAVISVILLE YARD	MUIE	0	0	NaN	NaN	
4097	2021-04-13	10:57	Tuesday	SPADINA STATION	SUAE	0	0	NaN	NaN	
4119	2021-04-13	22:00	Tuesday	YONGE- UNIVERSITY AND B	MUO	0	0	NaN	NaN	
4336	2021-04-19	23:00	Monday	SHEPPARD WEST TO LAWRE	MUO	0	0	NaN	NaN	
4748	2021-04-29	22:00	Thursday	YONGE UNIVERSITY SPADI	MUO	0	0	NaN	NaN	

	date	time	day	station	code	min_delay	min_gap	bound	line	vehi
5312	2021-05-15	05:05	Saturday	SPADINA BD STATION	MUNCA	0	0	NaN	NaN	
5448	2021-05-18	20:15	Tuesday	VAUGHAN MC STATION	MUWR	0	0	NaN	NaN	
5484	2021-05-19	18:11	Wednesday	Queen's Quay Station	PUMEL	0	0	NaN	NaN	
5642	2021-05-23	23:19	Sunday	ST ANDREW STATION	SUDP	0	0	NaN	NaN	
5685	2021-05-25	00:19	Tuesday	DUNDA WEST STATION	SUO	0	0	E	NaN	
6042	2021-06-02	22:28	Wednesday	WARDEN STATION	MUIRS	0	0	NaN	NaN	
6046	2021-06-02	00:56	Wednesday	BAY STATION	SUO	0	0	NaN	NaN	
6540	2021-06-14	22:43	Monday	YONGE BD STATION	MUIS	0	0	NaN	NaN	
6560	2021-06-15	07:15	Tuesday	SUBWAY OPERATIONS BUIL	PUMEL	0	0	NaN	NaN	

	date	time	day	station	code	min_delay	min_gap	bound	line	veh
7137	2021-06-28	01:03	Monday	COXWELL STATION	MUNCA	0	0	NaN	NaN	
7766	2021-07-14	03:51	Wednesday	TRANSIT CONTROL CENTRE	PUSO	0	0	NaN	NaN	
8889	2021-08-11	07:46	Wednesday	TRANSIT CONTROL	MUIE	0	0	NaN	NaN	
9628	2021-08-29	15:49	Sunday	YORKDALE STATION	SUPOL	0	0	NaN	NaN	
9629	2021-08-29	16:13	Sunday	YORK MILLS STATION	MUO	0	0	NaN	NaN	
9780	2021-09-01	20:35	Wednesday	MAIN STREET AND UNION	MUO	0	0	NaN	NaN	
9789	2021-09-01	22:14	Wednesday	UNION AND KENNEDY STAT	MUO	0	0	NaN	NaN	
10336	2021-09-13	17:20	Monday	MCBRIEN BUILDING	SUO	0	0	NaN	NaN	
10951	2021-09-26	15:50	Sunday	WILSON STATION	PUOPO	0	0	N	NaN	54
11223	2021-10-01	00:33	Friday	WELLESLEY STATION	SUDP	0	0	NaN	NaN	
12533	2021-10-28	14:18	Thursday	VICTORIA	MUIS	0	0	NaN	NaN	

	date	time	day	station	code	min_delay	min_gap	bound	line	vehi
				PARK STATION						
12826	2021-11-02	12:22	Tuesday	GREENWOOD SHOP	MUIE	0	0	NaN	NaN	
13007	2021-11-05	08:59	Friday	BLOOR STATION	MUIRS	0	0	S	NaN	
13080	2021-11-06	18:41	Saturday	KENNEDY BD STATION	MUO	0	0	NaN	NaN	
13273	2021-11-10	16:25	Wednesday	SUMMERHILL STATION	TUS	3	6	N	NaN	65
13402	2021-11-12	20:42	Friday	CLOSURES BUILDING	MUIE	0	0	NaN	NaN	
13410	2021-11-12	00:02	Friday	TRANSIT CONTROL	MUO	0	0	NaN	NaN	
14177	2021-11-25	21:14	Thursday	WILSON CARHOUSE	MUIE	0	0	NaN	NaN	
14371	2021-11-29	05:10	Monday	GO PROTOCOL	MUO	0	0	NaN	NaN	

	date	time	day	station	code	min_delay	min_gap	bound	line	vehi
14935	2021-12-08	06:00	Wednesday	TORONTO TRANSIT COMMIS	MUO	0	0	NaN	NaN	
14952	2021-12-08	13:58	Wednesday	KIPLING STATION	MUIS	0	0	NaN	NaN	
14967	2021-12-08	17:14	Wednesday	QUEEN'S PARK STATION	MUO	0	0	NaN	NaN	
15581	2021-12-19	00:42	Sunday	WILSON CARHOUSE	MUO	0	0	NaN	NaN	
15623	2021-12-20	16:23	Monday	YONGE- SHEPPARD (LINE 4	MUIRS	0	0	NaN	NaN	
16332	2021-12-31	14:34	Friday	GO PROTOCOL	MUO	0	0	NaN	NaN	

.loc[] also lets us access data by label, with row conditions first and column conditions second.

### Out[40]:

	time	station	line
495	15:22	FINCH WEST STATION	NaN
513	22:08	EGLINTON WEST STATION	NaN
1044	22:00	YONGE-UNIVERSITY AND B	NaN
1045	23:00	FINCH STATION	NaN
1362	01:45	LAWRENCE STATION	NaN

### query()

Alternatively, we can use the DataFrame query() method, which takes a filter condition as a string, and returns a DataFrame of records that met the condition. query() is slower than loc[], but it can be easier to read.

### Out[41]:

	date	time	day	station	code	min_delay	min_gap	bound	line	vehicle
495	2021-01-13	15:22	Wednesday	FINCH WEST STATION	MUSAN	3	6	S	NaN	5751
513	2021-01-13	22:08	Wednesday	EGLINTON WEST STATION	PUMEL	0	0	NaN	NaN	0
1044	2021-01-27	22:00	Wednesday	YONGE- UNIVERSITY AND B	MUO	0	0	NaN	NaN	0
1045	2021-01-27	23:00	Wednesday	FINCH STATION	MUO	0	0	NaN	NaN	0
1362	2021-02-04	01:45	Thursday	LAWRENCE STATION	TUSC	0	0	S	NaN	5596

### dropna()

In this case, the number of records without lines is relatively small. Most do not have delay durations. Some appear to be at rail yards, i.e. not on a rail line. For our analysis, we may drop them with the dropna() DataFrame method. We can drop rows missing lines by passing a subset.

```
In [42]: delays_w_reasons = delays_w_reasons.dropna(subset='line')
```

Filtering data with .loc[] and isin()

We can use .loc[] to create a delays DataFrame without the invalid lines. To to this, we first create a list of values to exclude, then pass the list to the Series isin() method. Finally, we negate the expression, and assign the output back to delays\_w\_reasons.

Note: The negation operator here is ~, not !. The and or operators are different as well: & and | respectively.

```
In [43]: # set up filter list
    filter_list = ['999', '36 FINCH WEST', '35 JANE', '52', '41 KEELE']
```

Replacing values with str.replace()

To standardize the YU/BD values, we can replace the less common ones. One way to do this is by selecting the line Series and using str.replace(), like below for "YUS".

Another is to assign "YU/BD" to values selected by .loc[]

#### Out[46]:

590	10Y	NGE/UNI	VEF	RS]	TY/I	BL(	OOR
852					YU	/	BD
1137					YU	/	BD
1628					YU	/	BD
1672					YU	/	BD
1700					YU	/	BD
6725					YU	/	BD
7469					YU	/	BD
8034					YU	&	BD
8301					YU	/	BD
8341					YU	/	BD
8463					YU	/	BD
9164					YU	/	BD
9541		,	YU	&	BD I	LI	NES
9839					YU	/	BD
10792					YU	/	BD
11119					YU	/	BD
11299					YU	/	BD
12128					YU	/	BD
15574					YU	/	BD
Name:	line,	dtype:	ol	ρje	ect		

```
Out[47]:
array(['YU', 'BD', 'SHP', 'SRT', 'YU/BD', 'SHEP'], dtype=object)
```

# Practice

- Select the date, time, and station for all delays on the Yonge-University ("YU") and Bloor-Danforth ("BD") lines.
- Create a new DataFrame, nonzero\_delays, containing delays with a min\_time greater than zero.

## Grouping

A core workflow in pandas is *split-apply-combine*:

- **splitting** data into groups
- **applying** a function to each group, such calculating group sums, standardizing data, or filtering out some groups
- **combining** the results into a data structure

This workflow starts by grouping data by calling the groupby() method. We'll pass in a column name or list of names to group by.

```
In [48]: line_groups = delays_w_reasons.groupby('line')
```

groupby() returns a grouped DataFrame that we can use to calculate groupwise statistics. The grouping column values become indexes, or row labels. **Note that this grouped DataFrame still references the original, so mutating one affects the other.** 

```
In [49]: # how many hours of delays did each line have in 2021?
    line_groups['hour_delay'].sum()
```

#### Out[49]:

```
line
BD 329.47
SHEP 0.00
SHP 28.43
SRT 57.82
YU 477.50
YU/BD 0.00
```

Name: hour\_delay, dtype: float64

We can group by more than one column by passing a list into groupby(). Data is grouped in the order of column names.

Chaining methods and unstack() ing

We can *chain* methods together for convenience and code readability. Here, we calculate the size() of each group, then unstack() the resulting Series by the first part of the row label, line. The tail() method is added to the end so that the output takes less screen space.

#### Out[51]:

line	BD	SHEP	SHP	SRT	YU	YU/BD
code_description						
Work Refusal	4.0	NaN	1.0	NaN	12.0	NaN
Work Vehicle	3.0	NaN	NaN	NaN	7.0	NaN
Work Zone Problems - Signals	4.0	NaN	4.0	NaN	5.0	NaN
Work Zone Problems - Track	12.0	NaN	NaN	NaN	29.0	NaN
Yard/Carhouse Related Problems	17.0	NaN	NaN	NaN	15.0	NaN

```
agg() regating
```

So far, we have applied one function at a time. The agg() DataFrame method lets us apply multiple functions on different columns at once.

agg() 's argument syntax is a little unusual. It follows this pattern:

#### Out[52]:

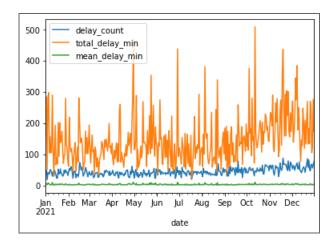
	delay_count	total_delay_min	mean_delay_min
date			
2021-01-01	36	159	4.416667
2021-01-02	49	284	5.795918
2021-01-03	19	51	2.684211
2021-01-04	41	284	6.926829
2021-01-05	40	298	7.450000

## Plotting

The summary table we just generated can be easily plotted within pandas. Since the index contains dates, pandas automatically knows to plot values as time series data, with the dates in the x-axis -- we just have to call the plot() method.

In [53]: delay\_summary.plot()

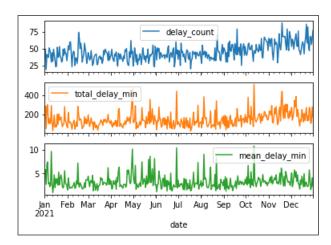
Out[53]:
 <AxesSubplot:xlabel='date'>



### To create a separate plot for each column, we can specify subplots=True

```
In [54]: delay_summary.plot(subplots=True)
```

#### Out[54]:

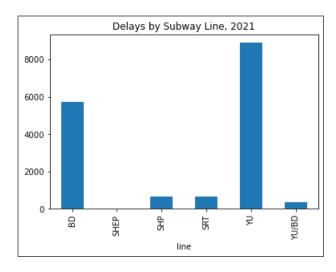


We can plot other aggregations too. Below, we use  $line\_groups$  and calculate the size of each group, i.e., the number of delays reported on each line. Then we plot the data, telling pandas that the plot kind should be a bar graph, with TTC lines should in the x-axis. We also pass in a title.

```
In [55]: (line_groups
          .size()
.plot(x='line',
          kind='bar',
          title='Delays by Subway Line, 2021'))
```

#### Out[55]:

<AxesSubplot:title={'center':'Delays by Subway Line, 2021'}, xlabel='line'>



# Practice

- Sum up and plot the total delay time, in hours, by line.
- Re-create delay\_summary, but group by date and line this time. Add an additional aggregation column.
  - Experiment with the unstack() and reset\_index() methods. Try passing in 0 or 1 as an argument. How does delay\_summary change?

# Writing to file

## **Exporting DataFrames**

DataFrames have to\_[file format]() methods, analogous to pandas read functions. The counterpart to pd.read\_csv(), for instance, is DataFrame.to\_csv(). The export methods generally take a file path to save to as their first argument. Additional arguments vary a bit by export format, but index is a common useful one. It takes a boolean of whether to write the index to file -- set it to False if the index is the numbered default.

Two additional useful parameters in DataFrame.to\_csv() and DataFrame.to\_excel() are na\_rep, which takes a string to use for null values, and columns, which lets us write out a subset of columns.

```
In [57]: # write delays_w_reasons to an Excel file
    delays_w_reasons.to_excel('ttc_subway_delays_w_reasons.xlsx', index=False)
```

# More wrangling

### Neighbourhood Profiles

The bike theft data includes neighbourhood identifiers. These neighbourhoods are designated by City of Toronto, which publishes neighbourhood demographic profiles. Let's get this data to start investigating if neighbourhoods with more bike theft reports simply have higher populations. In the process, we will reinforce what we learned so far. We will also learn about two last ways to reshape data: melt(), which rearranges data from a wide format to a long format; and pivot(), which reorganizes data based on index and column values.

# Getting data

Let's load the neighbourhood data and explore it.

```
In [58]: profiles = pd.read_csv('../data/neighbourhood-profiles-2016-140-model.csv')
In [59]: profiles.shape
Out[59]:
  (2383, 146)
```

The neighbourhood profiles are in an unusual format. Neighbourhood names are in the columns, while attribute fields are rows, and there are thousands of attributes.

In [60]: profiles.head()

Out[60]:

	_id	Category	Торіс	Data Source	Characteristic	City of Toronto	Agincourt North	Sout
0	1	Neighbourhood Information	Neighbourhood Information	City of Toronto	Neighbourhood Number	NaN	129	
1	2	Neighbourhood Information	Neighbourhood Information	City of Toronto	TSNS2020 Designation	NaN	No Designation	Des
2	3	Population	Population and dwellings	Census Profile 98-316- X2016001	Population, 2016	2,731,571	29,113	
3	4	Population	Population and dwellings	Census Profile 98-316- X2016001	Population, 2011	2,615,060	30,279	
4	5	Population	Population and dwellings	Census Profile 98-316- X2016001	Population Change 2011-2016	4.50%	-3.90%	

Because of the layout and formatting characters, all of the numeric values have been read in as text data. It also looks like the characteristics are not unique.

## Removing extra whitespace

The characteristic values contain extra whitespace. Let's remove the whitespace up with str.strip().

```
In [63]: # the whitespace is easier to see in a list than a Series
        list(profiles['Characteristic'][95:100])
Out[63]:
      Female parent',
      Male parent',
 'Couple census families in private households',
    Couples with children',
      1 child']
In [64]: profiles['Characteristic'] = profiles['Characteristic'].str.strip()
# get the first 10 characteristics
list(profiles['Characteristic'][95:100])
Out[64]:
['Female parent',
  'Male parent',
 'Couple census families in private households',
  'Couples with children',
 '1 child']
```

## Subsetting data

1651 characteristics is still a lot. Let's check out the categories to understand the areas covered.

"Journey to work" sounds relevant. We can use .loc[] to get the rows in that category, then select the Topic column and get its unique values.

The "Population and dwellings" topic we saw in the DataFrame head looked promising as well. Let's check out the Characteristics in that topic.

Now that we know what topics we're interested in, let's create a subset DataFrame limited to them. We'll use the copy() DataFrame method to leave the original data untouched.

```
In [68]: topics = ['Neighbourhood Information', 'Population and dwellings', 'Main mode of commuting']
# make sure it's an independent copy
profiles_subset = profiles.copy()
# get just the topics we're interested in
profiles_subset = profiles_subset.loc[profiles['Topic'].isin(topics)]
profiles_subset.shape
```

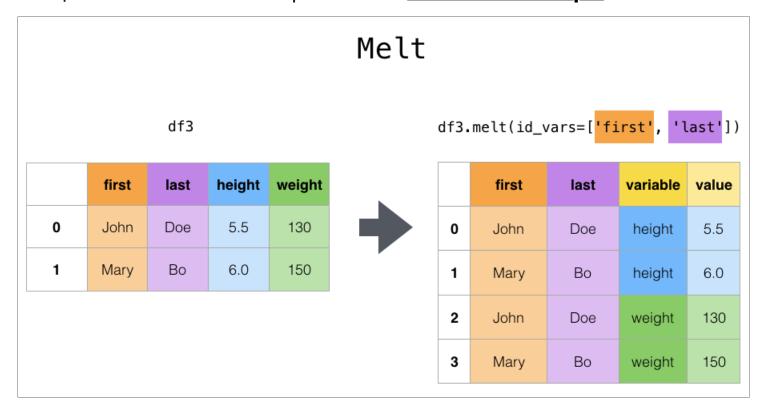
```
Out[68]: (16, 146)
```

# Reshaping data with melt()

Now we're ready to reshape our data. We can <code>drop()</code> the ID, data source, and category columns now.

To melt() a DataFrame, we specify id\_vars -- the columns to keep as identifiers. All other columns are 'melted' into a new variable column. The values at DataFrame[id\_vars, variable\_col] move into a value column. We can change the names of the variable and value columns with the var\_name and value\_name arguments.

# The pandas documentation provides an **illustrative example**:



Let's melt() the profiles subset. We'll keep Topic and Characteristic as our id\_vars. This will melt the neighbourhood names into the variable column, which we'll rename Neighbourhood.

### Out[69]:

	Торіс	Characteristic	Neighbourhood	value
0	Neighbourhood Information	Neighbourhood Number	City of Toronto	NaN
1	Neighbourhood Information	TSNS2020 Designation	City of Toronto	NaN
2	Population and dwellings	Population, 2016	City of Toronto	2,731,571
3	Population and dwellings	Population, 2011	City of Toronto	2,615,060
4	Population and dwellings	Population Change 2011-2016	City of Toronto	4.50%

# Reshaping data with pivot()

The profile data is looking much closer to what we want! The next step is to make the Topic/Characteristic the column header, pivot() ing the values. To do this, we specify the column(s) to use as the index, or row labels; the column(s) whose values we should use as column names, and which column our values come from.

Pivoting on two columns creates a multi-level column header, so we then drop the top Topic level with droplevel(). Finally, we reset\_index() to make neighbourhood names a regular column.

Out[70]:

									Privat
			Neighbourhood	TSNS2020	Population,	Population	<b>Population</b>	Total	dwelling
Char	acteristic	Neighbourhood	Number	Designation	2016	2011	Change	•	occupie
							2011-2016	dwellings	by usua residen

0	Agincourt North	129	No Designation	29,113	30,279	-3.90%	9,371	9,12
1	Agincourt South-	128	No Designation	23,757	21,988	8.00%	8,535	8,13

N	1alvern
	West

2	Alderwood	20	No Designation	12,054	11,904	1.30%	4,732	4,61
3	Annex	95	No Designation	30,526	29,177	4.60%	18,109	15,934
4	Banbury- Don Mills	42	No Designation	27,695	26,918	2.90%	12,473	12,12

# Renaming all columns

Much better! These column names could be shorter, though. Let's rename them to be easier to work with. We could use the rename() DataFrame method, passing in a dictionary of old and new names. Since there isn't an easy renaming function, and some of the current names are very long, we will instead reassign a list of new names to the columns attribute of our DataFrame.

```
In [71]: # rename all columns
        neighbourhoods.columns = ['neighbourhood',
                           'n_id',
                           'designation',
                           'pop_2016',
                           'pop 2011',
                           'pop change',
                           'private dwellings',
                           'occupied_dwllings',
                           'pop dens',
                           'area',
                           'total commuters',
                           'drive',
                           'car_passenger',
                           'transit',
                           'walk',
                           'bike',
                           'other']
neighbourhoods.columns
```

#### Out[71]:

# Replacing values in multiple columns

All of the values in our neighbourhood data are text right now. Part of the problem is that numbers contain characters like commas and percentage signs. We can remove these from everywhere in our data with the DataFrame replace() method, which takes a string to look for and a replacement string. Normally, replace() looks for a perfect, full-string match. Since we're only looking for a substring match, we set regex=True.

### Out[72]:

	neighbourhood	n_id	designation	pop_2016	pop_2011	pop_change	private_dwellings	occupied_dwllings	pop_c
0	Agincourt North	129	No Designation	29113	30279	-3.90	9371	9120	39
1	Agincourt South- Malvern West	128	No Designation	23757	21988	8.00	8535	8136	30

# apply() ing a function to multiple columns

Now the numbers look like numbers, but they are still strings. We can convert them with pd.to\_numeric(), which takes a Series and returns it as the most appropriate numeric data type. Doing this for columns one-by-one would be tedious. Instead, we can use the apply() DataFrame method to run a function on every column in a DataFrame. apply() takes the name of the function to apply and any arguments needed to run that function. We only want to convert from pop\_2016 onwards, so we'll use .loc[] to select the correct columns.

### Out[73]:

	neighbourhood	n_id	designation	pop_2016	pop_2011	pop_change	private_dwellings	occupied_dwllings	pop_c
0	Agincourt North	129	No Designation	29113	30279	-3.9	9371	9120	39
1	Agincourt South- Malvern West	128	No Designation	23757	21988	8.0	8535	8136	30
2	Alderwood	20	No Designation	12054	11904	1.3	4732	4616	24
3	Annex	95	No Designation	30526	29177	4.6	18109	15934	108
4	Banbury- Don Mills	42	No Designation	27695	26918	2.9	12473	12124	27

# 

### Out[74]:

object object object
int64
int64
float64
int64
int64
int64
float64
int64

# Calculating more columns

Let's fix the population change column and calculate the percentage of commuters who bike.

### Out[75]:

	neighbourhood	n_id	designation	pop_2016	pop_2011	pop_change	private_dwellings	occupied_dwllings	pop_c
0	Agincourt North	129	No Designation	29113	30279	-0.039	9371	9120	39
1	Agincourt South- Malvern West	128	No Designation	23757	21988	0.080	8535	8136	30
2	Alderwood	20	No Designation	12054	11904	0.013	4732	4616	24
3	Annex	95	No Designation	30526	29177	0.046	18109	15934	108
4	Banbury- Don Mills	42	No Designation	27695	26918	0.029	12473	12124	27

merge() ing

The profile are now ready to merge into the bike thefts data!

Out[76]:							
	_id	objectid	event_unique_id	primary_offence	occurrence_date	occurrence_year	occurrence_month
0	1	17744	GO-20179016397	THEFT UNDER	2017-10-03	2017	October
1	2	17759	GO-20172033056	THEFT UNDER - BICYCLE	2017-11-08	2017	November
2	3	17906	GO-20189030822	THEFT UNDER - BICYCLE	2018-09-14	2018	September
3	4	17962	GO-2015804467	THEFT UNDER	2015-05-07	2015	May

	_id	objectid	event_unique_id	primary_offence	occurrence_date	occurrence_year	occurrence_month
4	5	17963	GO-20159002781	THEFT UNDER	2015-05-16	2015	May
•••	•••	•••		•••	•••	•••	
25564	25565	9361	GO-2015543181	MISCHIEF UNDER	2015-04-01	2015	April
25565	25566	11318	GO-20169004589	THEFT UNDER	2016-05-16	2016	May
25566	25567	11462	GO-20169005434	THEFT UNDER	2016-06-04	2016	June
25567	25568	11695	GO-20161170896	THEFT UNDER	2016-07-04	2016	July

				THEFT			
25568	25569	11883	GO-20169007653	UNDER -	2016-07-22	2016	July
				BICYCLE			

event\_unique\_id primary\_offence occurrence\_date occurrence\_year occurrence\_month

25569 rows × 51 columns

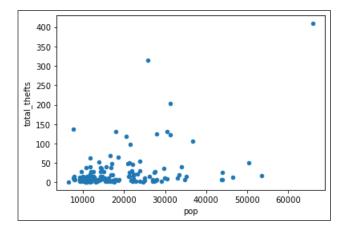
\_id objectid

# Grouping and plotting

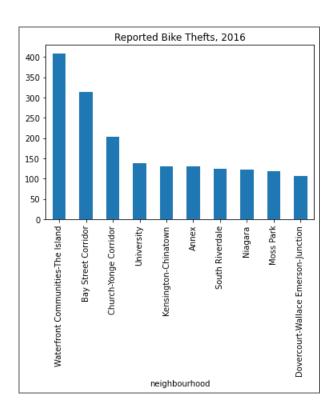
With the datasets joned, we can aggregate and plot the data. We can start using statistical methods, like corr() to check for relationships between variables as well.

### Out[81]:

<AxesSubplot:xlabel='pop', ylabel='total\_thefts'>



Out[82]:
 <AxesSubplot:title={'center':'Reported Bike Thefts, 2016'}, xlabel='neighbourhood'>



```
In [83]: # a quick correlation check
          (thefts_2016_grouped
          agg(total_thefts=('_id', 'count'),
          pop=('pop_2016', 'median'),
          dens=('pop_dens', 'median'),
          pct_bike=('pct_bike', 'mean'))
          corr('spearman'))
```

### Out[83]:

	total_thefts	рор	dens	pct_bike
total_thefts	1.000000	0.267761	0.485556	0.650101
рор	0.267761	1.000000	0.020082	-0.218934
dens	0.485556	0.020082	1.000000	0.603434
pct_bike	0.650101	-0.218934	0.603434	1.000000

# References

## Programming

- pandas development team. API reference. <a href="https://pandas.pydata.org/pandas-docs/stable/reference/index.html">https://pandas.pydata.org/pandas-docs/stable/reference/index.html</a>
- pandas development team. *User guide*. <a href="https://pandas.pydata.org/pandas-docs/stable">https://pandas.pydata.org/pandas-docs/stable</a> /user\_guide/index.html
- Python strftime cheatsheet. <a href="https://strftime.org/">https://strftime.org/</a>

### **Data Sources**

- Open Data Toronto. Neighbourhood Profiles. <a href="https://open.toronto.ca/dataset/">https://open.toronto.ca/dataset/</a>
   /neighbourhood-profiles/
- Open Data Toronto. TTC Subway Delay Data. <a href="https://open.toronto.ca/dataset/ttc-subway-delay-data/">https://open.toronto.ca/dataset/ttc-subway-delay-data/</a>
- Open Data Toronto. Bicyle Thefts. <a href="https://open.toronto.ca/dataset/bicycle-thefts/">https://open.toronto.ca/dataset/bicycle-thefts/</a>